

# Education Data Science Summit

Ryan Estrellado

4/30/2021

This document is meant to accompany my talk at the Education Data Science Summit. It uses publicly available data. I hope it helps you!

Ryan Estrellado

[linktr.ee/ry\\_estrellado](http://linktr.ee/ry_estrellado)

## What is an R Markdown Document?

The authors of **R Markdown: The Definitive Guide** write that R Markdown empowers data scientists to “interweave narratives with code in a document . . .”

Think of it as a Word document that holds not only story, but also code and graphs.

## Things to Watch For

Look for how writing, code, and plots appear in a single document.

Look for the parts of good data analysis using code:

- Identify the Question
- Import the Data
- Prepare the Data
- Pick a method to answer the question
- Visualize the data
- Analyze the Data
- Report the findings

## Load Packages

Packages are add-ons to R that are built by software developers and communit members. They extend the capability of base R.

```
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.0 --

## v ggplot2 3.3.3      v purrr   0.3.4
## v tibble  3.0.6      v dplyr  1.0.4
## v tidyr   1.1.2      v stringr 1.4.0
## v readr   1.4.0      v forcats 0.5.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```
library(here)
```

```
## here() starts at /Users/restrellado/OneDrive - San Diego County Superintendent of Schools/2017-2018/
```

```
library(janitor)
```

```
##
```

```
## Attaching package: 'janitor'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      chisq.test, fisher.test
```

## Identify the Question

How many schools do districts have on average?

## Import Data

Before you work with data in R, you have to import the data.

You can get this data from the California Department of Education's website:

```
enroll <- read_tsv(here::here("19-20-enrollment.txt"))
```

```
##
```

```
## -- Column specification -----
```

```
## cols(
```

```
##   .default = col_double(),
```

```
##   CDS_CODE = col_character(),
```

```
##   COUNTY = col_character(),
```

```
##   DISTRICT = col_character(),
```

```
##   SCHOOL = col_character(),
```

```
##   GENDER = col_character()
```

```
## )
```

```
## i Use `spec()` for the full column specifications.
```

Here's one way working in R is different from working in a spreadsheet: you aren't dragging and dropping things around in a spreadsheet. Instead, you're using code to submit instructions to R.

Sometimes you just want to see the dataset. Here's a way to do that:

```
# See the data using View()
```

```
# View(enroll)
```

Question: Let's look at a few variables. Do these variables contain unique values? In other words, does Castle Rock school appear only once?

## Prepare Data

Most publicly available datasets aren't prepared for the analysis you want to do. You have to prepare the dataset before you work with it.

- We want the average number of schools in a district
- To get that, we need to be able to count the number of schools in each district
- And to do that, we need each school in each district to only appear once
- Last, we only want districts in San Diego county

```
sd <- enroll %>%
  filter(COUNTY == "San Diego") %>%
  distinct(DISTRICT, SCHOOL) # Each school should only appear once

#View(sd)
```

Question: How many schools does the Alpine district have?

Count the schools in each district:

```
school_count <- sd %>%
  count(DISTRICT, sort = TRUE)

school_count
```

```
## # A tibble: 50 x 2
##   DISTRICT      n
##   <chr>      <int>
## 1 San Diego Unified    223
## 2 Chula Vista Elementary    50
## 3 Poway Unified        40
## 4 Vista Unified        34
## 5 Sweetwater Union High    32
## 6 Cajon Valley Union     31
## 7 Escondido Union       26
## 8 Oceanside Unified       26
## 9 La Mesa-Spring Valley    25
## 10 San Marcos Unified     22
## # ... with 40 more rows
```

## Pick a Method to Answer the Question

Compute the average number of schools in a district. Compute the standard deviation of the number of schools in a district.

## Visualize Data

Make the dataset easier to visualize by filtering for the top ten:

```
top_10 <- school_count %>%
  filter(min_rank(desc(n)) < 10)

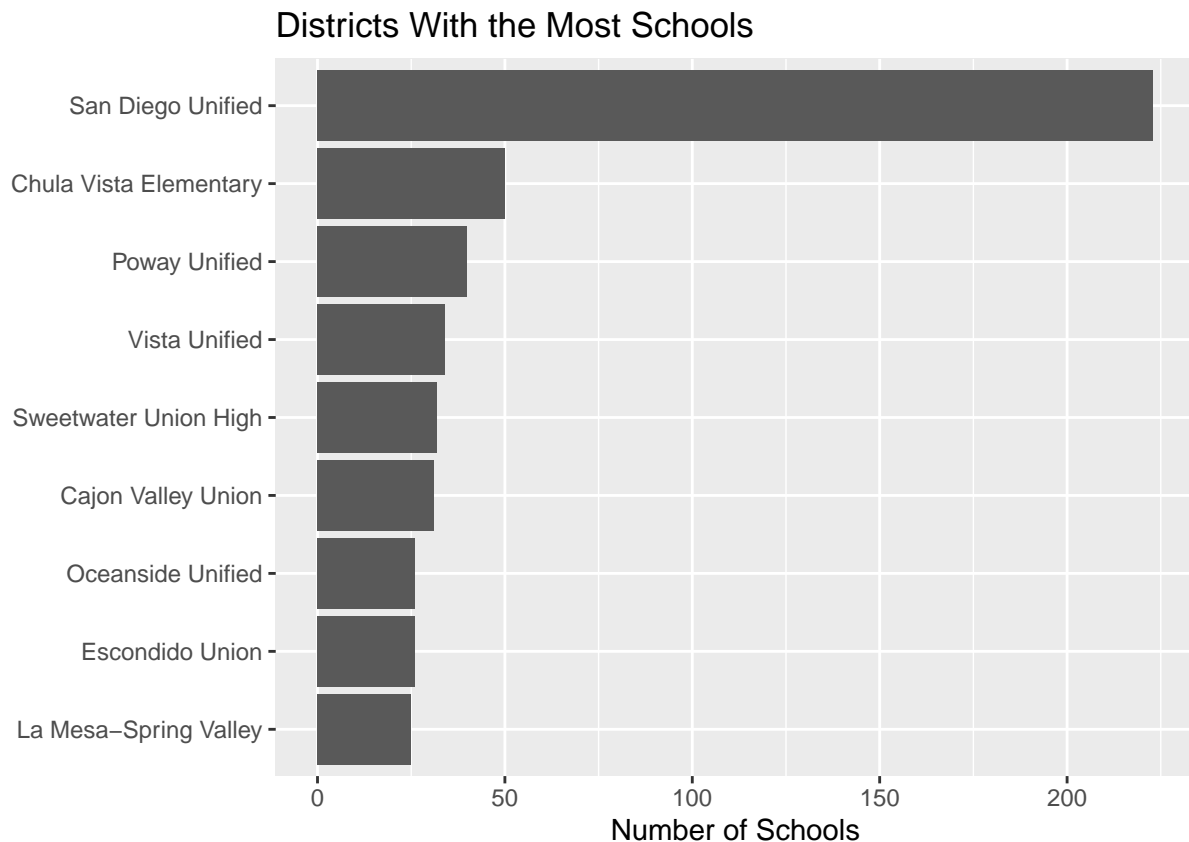
top_10
```

```
## # A tibble: 9 x 2
##   DISTRICT      n
##   <chr>      <int>
## 1 San Diego Unified    223
## 2 Chula Vista Elementary    50
## 3 Poway Unified        40
## 4 Vista Unified        34
## 5 Sweetwater Union High    32
## 6 Cajon Valley Union     31
## 7 Escondido Union       26
## 8 Oceanside Unified       26
```

```
## 9 La Mesa-Spring Valley      25
```

Visualize the top ten districts by number of schools. As I do this, watch how I talk and build the visualization layer by layer based on what I want to see.

```
ggplot(data = top_10, aes(x = reorder(DISTRICT, n), y = n)) +  
  geom_bar(stat = "identity") +  
  coord_flip() +  
  labs(title = "Districts With the Most Schools",  
        x = "",  
        y = "Number of Schools")
```



## Analyze the Data

Refresh our memory of the dataset:

```
#View(school_count)
```

. . . and let's compute the mean and standard deviation.

```
school_count %>%  
  summarise(avg_schools = mean(n), standard_dev = sd(n), median_schools = median(n))
```

```
## # A tibble: 1 x 3  
##   avg_schools standard_dev median_schools  
##       <dbl>       <dbl>         <dbl>  
## 1      15.9        31.9           9
```

## Report Your Findings

Districts in San Diego have about 16 schools in them, on average.

We should interpret this with caution, because there's a lot of variation in school counts in San Diego County. San Diego Unified has far and away the highest school count at 223 schools—almost. For the more statistics savvy, that's 7 times the standard deviation—a very large difference from the county average.